



ECE 57000

Basics of (supervised) Machine Learning

Chaoyue Liu

Fall 2024

Overview

First half of the semester: Supervised Learning

- Basics of machine Learning
- Linear model, linear regression, logistic regression
- Gradient descent, SGD
- Other non-neural networks models: kernel model, SVM (time-permitting)
- Neural network (fully-connected)
- Computer vision & Convolutional NN (CNN)
- Natural Language Processing (NLP) & Attention models

Second half of the semester: Unsupervised Learning

Supervised Machine Learning

- Data

- Dataset: a set of input-label pairs. “Supervised” because of existence of labels.
- Example: ImageNet -- more than **1 million** pairs of (image, label).



“goldfish”

terminology

input

label

notation

\mathbf{x}

y



An input can be: a vector (a list of numbers),
an image, a paragraph of text, ...

A label can be: real-valued (a number),
categorical (an element of a finite set)

Regression

If the label y is numeric (a real valued number), then the problem is known as regression.

- Example 1: given inputs (area, age, # of bedrooms, location, ...), predict/estimate the house's price.
- Example 2:



NOTE: Input x does not have to be numeric. Only the label y must be numeric.

Classification

If output is categorical,
then the problem is known as classification.

- Example 1: given height x , predict “male” ($y=0$) or “female” ($y=1$)
- Example 2: given salary x_1 and mortgage payment x_2 , predict defaulting on loan (“yes” or “no”)
- Example 3: given images, predict “cat” or “dog”

predicted: cat



predicted: cat



predicted: dog



predicted: cat



predicted: cat



predicted: dog



Question


MNIST:

- inputs: images with hand-written digits
- labels: digits from 0 to 9

Q: regression or classification?

A: classification.

- Predicting one digit: one element of a finite set
 - e.g., “5.5” makes no sense
- No ordering among digits
 - “4” is not closer to “5” than “1”



Supervised Learning

- Data
- Model

Build
the model

“machine”

Examples:
Neural network,
Kernel model,
SVM,
Linear model

...

Train
the model

“learning”

Use **training data**
and **algorithm**
to update the
model

Test/evaluate
the model


Use **test/dev data**
to evaluate the
learned model(s)

Now, let's use a simple model to demonstrate this procedure

k -NN: k -Nearest Neighbors

- a very simple and intuitive supervised learning model
- Intuition: data with similar inputs tend to have similar or the same labels

1. Choose k , $k=1,3,5\dots$ (for binary)
2. Find the k nearest neighbors in the training dataset.
3. Select most common class as the predicted label.



k -NN

Input: Test point x_0 , training data $D = \{(x_i, y_i)\}_{i=1}^n$

Output: Predicted class y_0

1. Compute distance to all training points:

$$d_i = \text{distance}(x_0, x_i), \forall i$$

2. Sort distances where π is a permutation:
(e.g., $\pi(1)$ is the index of the closest point)

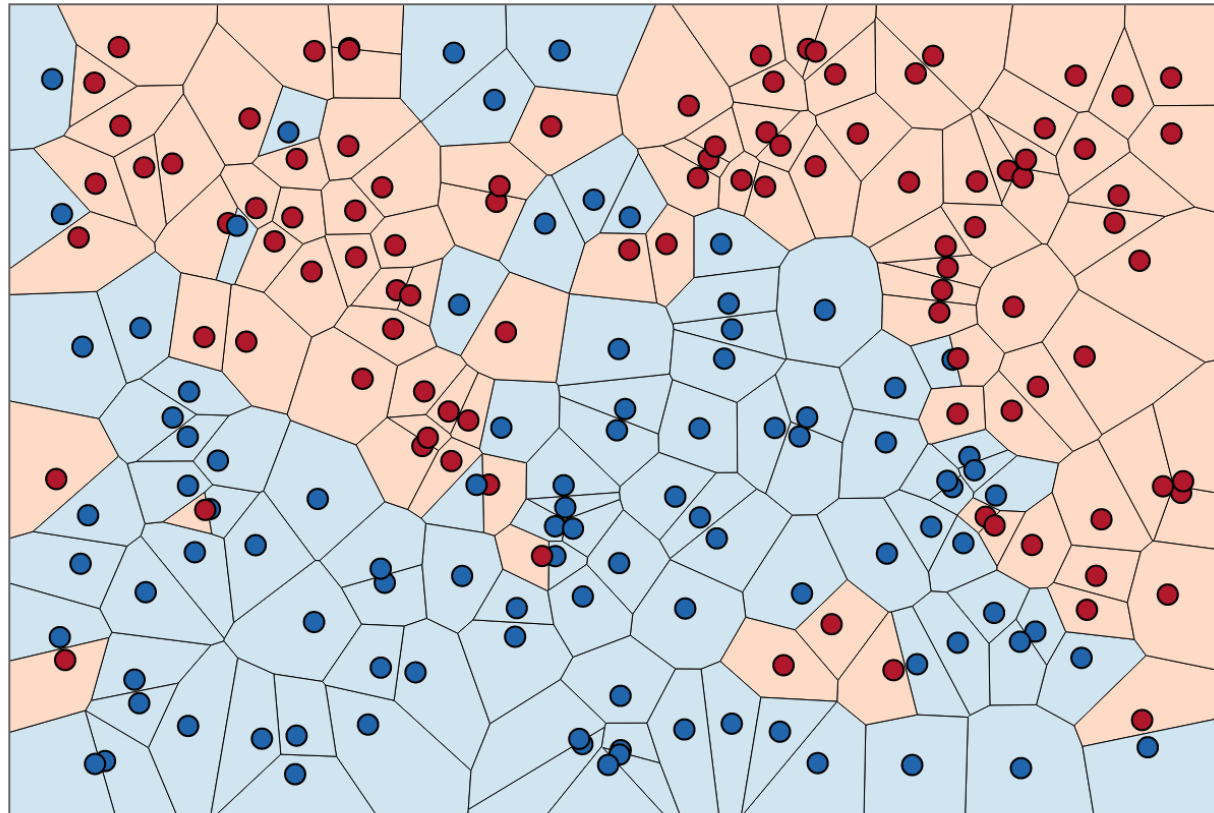
$$d_{\pi(1)} \leq d_{\pi(2)} \leq d_{\pi(3)} \leq \cdots \leq d_{\pi(n)}$$

3. Return the most common class of the top k (vote)

$$y_0 = \text{mode}(\{y_{\pi(1)}, y_{\pi(2)}, y_{\pi(3)}, \cdots, y_{\pi(k)}\})$$


1-NN ($k=1$)

1-NN partitions the space into Voronoi cells based on the training data




The k -NN boundary gets “smoother” as k increases

1-nearest neighbours



20-nearest neighbours



Now, let's go back to the procedure

Build
the model

Train
the model

Test/evaluate
the model

k -NN

Choose the k -NN
model and
Set k

no need to train
for k -NN
(a key part for
other models)

We talk about
this next...

Test/evaluate the model

Goal: make sure the model works well on **unseen** data, so that we can safely deploy the model

Q: can we use the same training data to test the model?

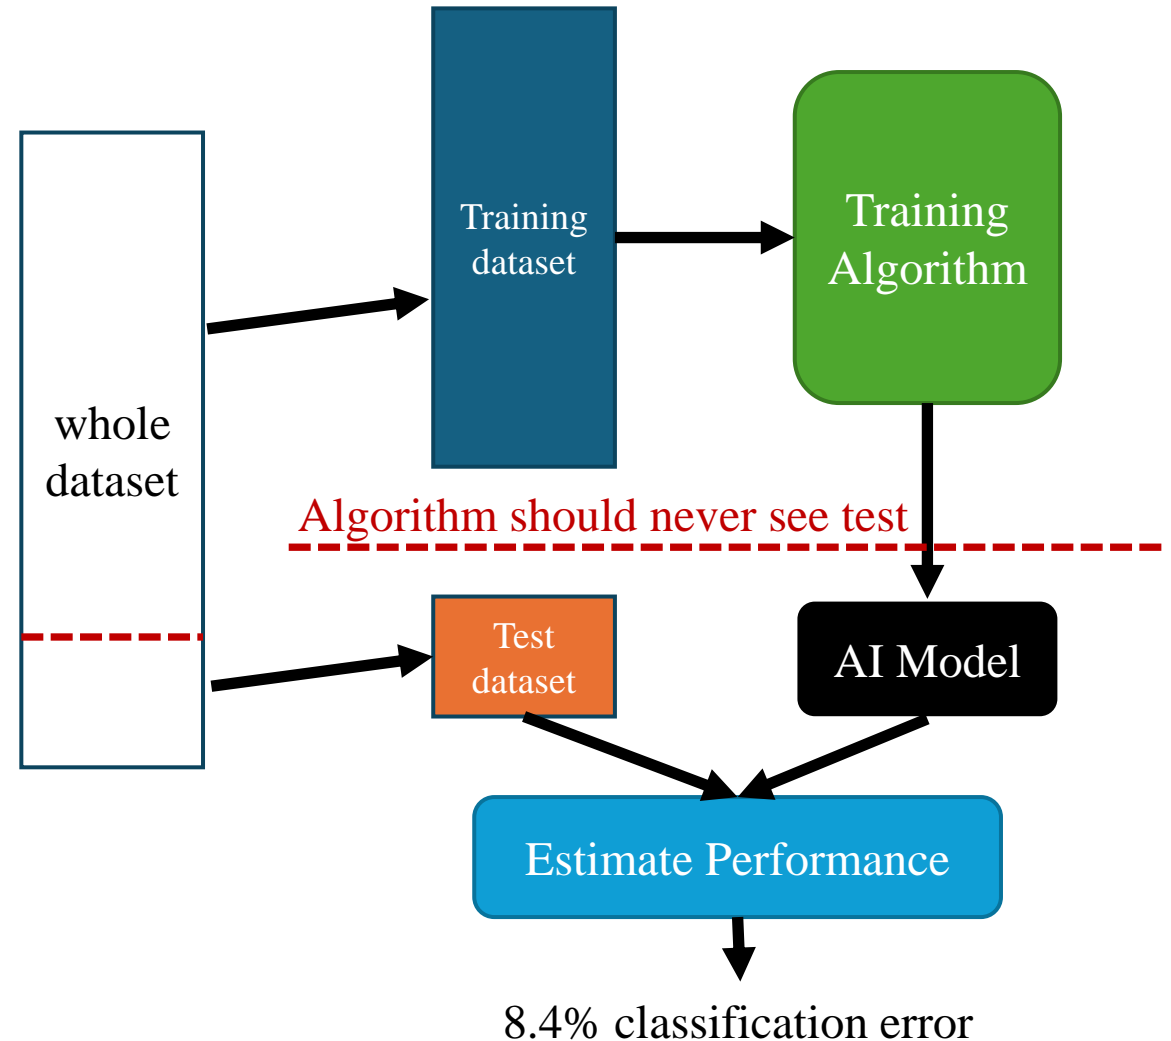
- If we train and evaluate on the same dataset, the model may only **memorize the training data** and **not generalize** well.
- Analogy to class
 - Training dataset is like homework, sample problems, and sample exams
 - Test dataset is like the real exam

But, we don't have the unseen data...

- estimate by splitting the known dataset

- Split the dataset

1. The training dataset is used to train the model only
2. The test dataset (or held-out dataset) is used to estimate generalization error.



k -NN: test performance

Input: Test dataset $D_{test} = \{(x_i, y_i)\}_{i=1}^m$, training dataset D_{train}

Output: generalization accuracy a

correct = 0

for each (x, y) in D_{test} :

 prediction $\tilde{y} = kNN(x, D_{train})$

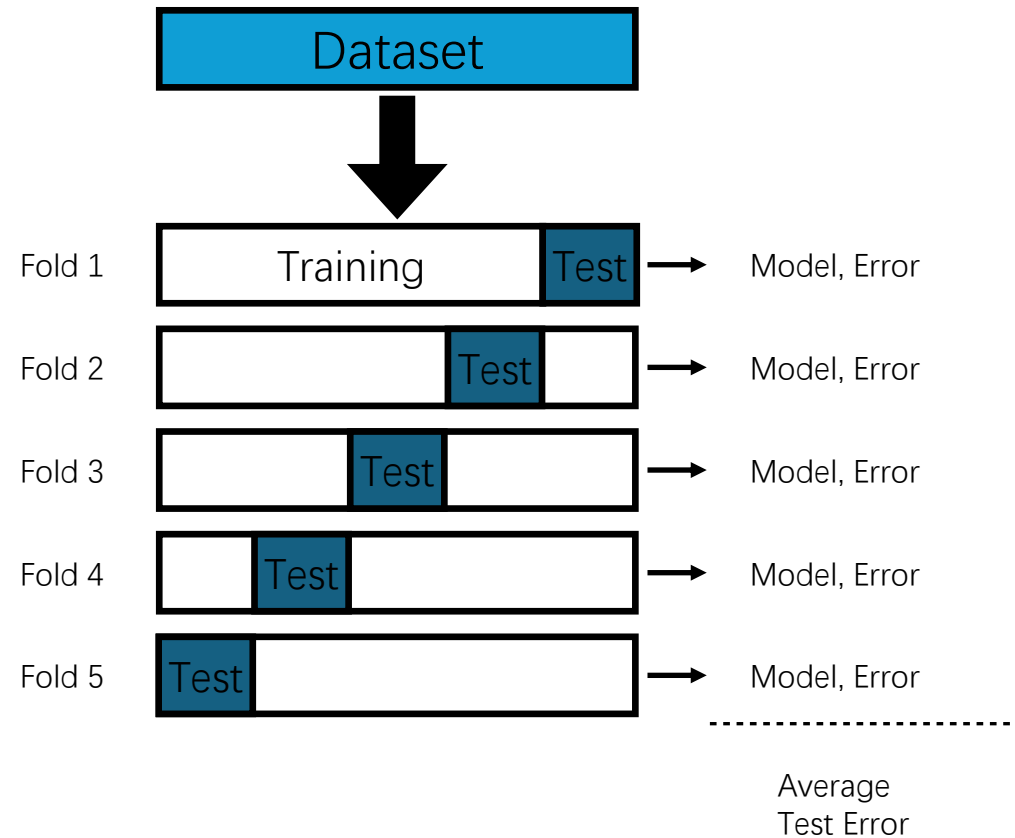
 if $\tilde{y}=y$, then correct = correct +1

generalization **accuracy** $a = \text{correct}/m$

generalization **error** $e = 1-a$.

Cross-validation (CV) generalizes the simple train/test split to M disjoint splits (effectively reusing data)

- Repeat the split process M times
 - Fit new model on train
 - Evaluate model on test
- Note: M models are fitted throughout process
- Final error estimate is average over all folds

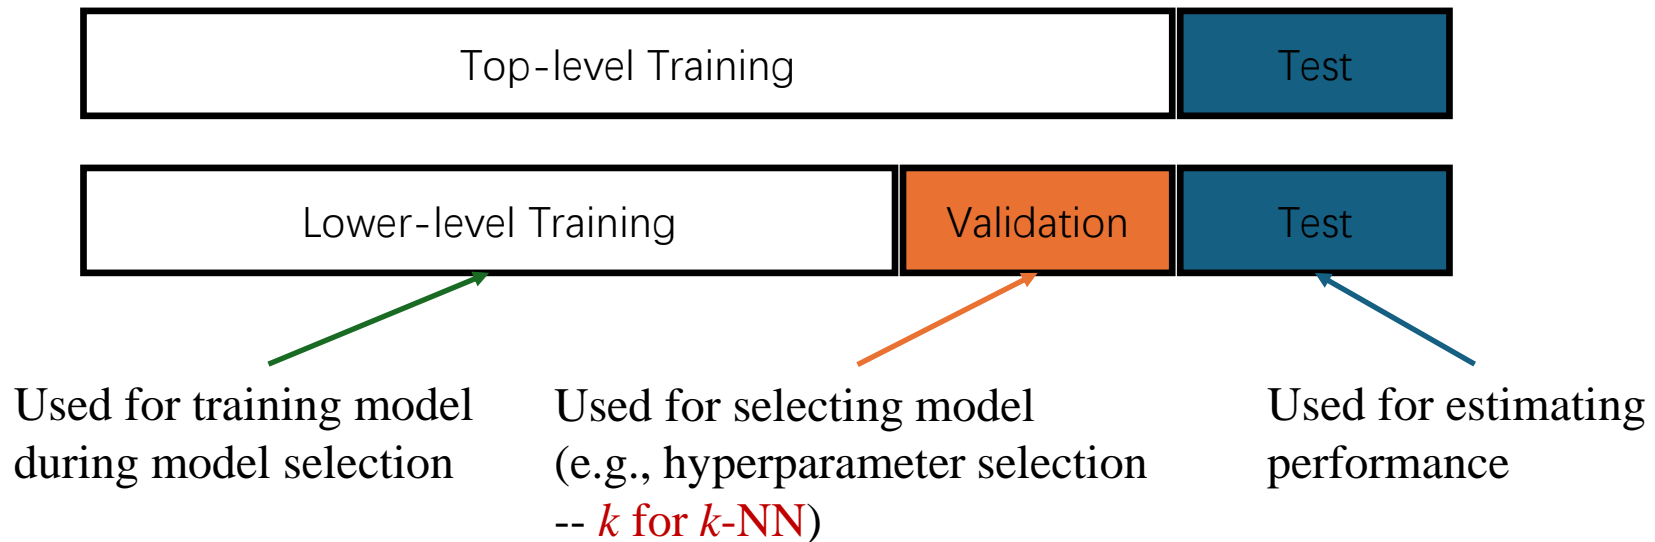


$M = 3, M = 5, M = 10$ are common

Cross-validation (CV) generalizes the simple train/test split to M disjoint splits


- Cross-validation is suitable for small dataset and small models
 - effectively reusing data, beneficial when dataset is small
 - requires running of M models, which is computationally expensive when model is large (e.g., deep neural networks).

But what if we want to select a model AND estimate the model's performance?



k -NN can be used for regression problems (predict continuous values)

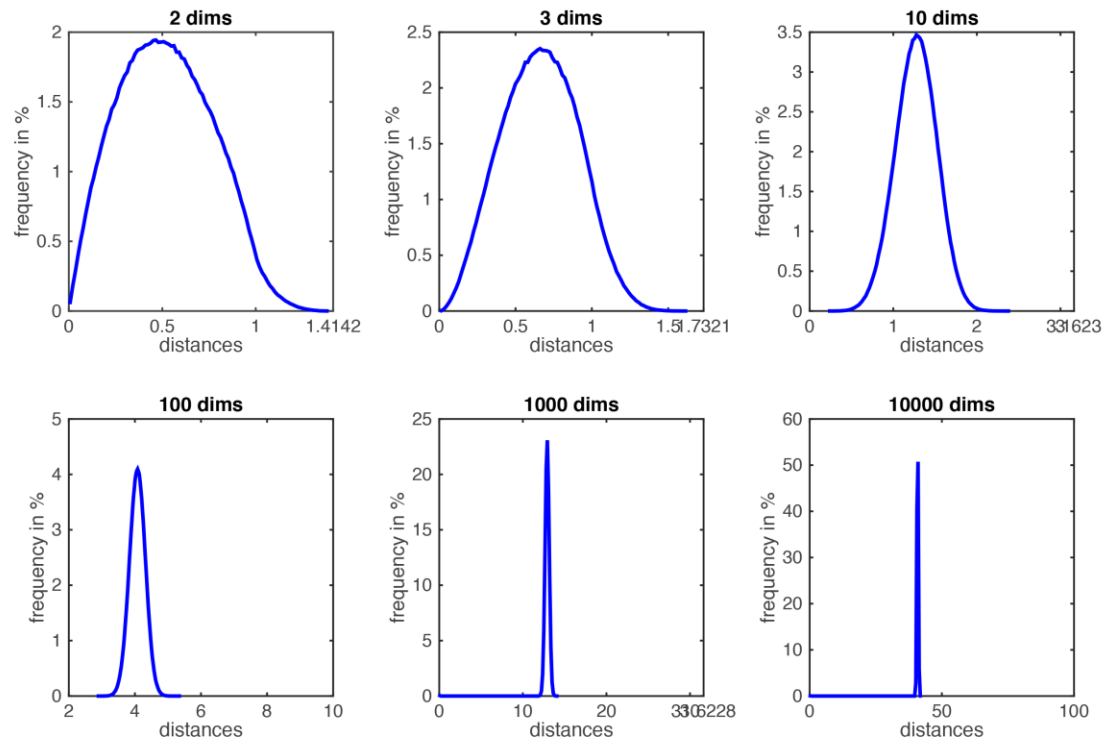
1. Find k nearest neighbors
2. Predict average of k nearest neighbors
(possibly weighted by distance)



Curse of dimensionality:

The performance of k -NN degrade significantly in high dimensions.

- The distances between **any two points** in high dimensions is nearly the same



Distance
between two
random points
concentrate
around a single
value

Related reading and source for KNN curse of dimensionality illustrations

- https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecturenote02_kNN.html