

```
In [1]: import torch
import numpy as np
import torch.nn as nn
import matplotlib.pyplot as plt
from torchvision.datasets import MNIST
from torch.utils.data import DataLoader
import torchvision.transforms as transforms
from mpl_toolkits.axes_grid1 import ImageGrid
from torchvision.utils import save_image, make_grid
```

```
In [2]: # create a transform to apply to each datapoint
transform = transforms.Compose([transforms.ToTensor()])

# download the MNIST datasets
path = '~/datasets'
train_dataset = MNIST(path, transform=transform, download=True)
test_dataset = MNIST(path, transform=transform, download=True)

# create train and test dataloaders
batch_size = 100
train_loader = DataLoader(dataset=train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=batch_size, shuffle=False)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
In [3]: class VAE(nn.Module):

    def __init__(self, input_dim=784, hidden_dim=400, latent_dim=200, device=device):
        super(VAE, self).__init__()

        # encoder
        self.encoder = nn.Sequential(
            nn.Linear(input_dim, hidden_dim),
            nn.LeakyReLU(0.2),
            nn.Linear(hidden_dim, latent_dim),
            nn.LeakyReLU(0.2)
        )

        # latent mean and variance
        self.mean_layer = nn.Linear(latent_dim, 2)
        self.logvar_layer = nn.Linear(latent_dim, 2)

        # decoder
        self.decoder = nn.Sequential(
            nn.Linear(2, latent_dim),
            nn.LeakyReLU(0.2),
            nn.Linear(latent_dim, hidden_dim),
            nn.LeakyReLU(0.2),
            nn.Linear(hidden_dim, input_dim)
        )

    def encode(self, x):
        x = self.encoder(x)
        mean, logvar = self.mean_layer(x), self.logvar_layer(x)
        return mean, logvar

    def reparameterization(self, mean, logvar):
        epsilon = torch.randn_like(logvar).to(device)
        z = mean + (logvar/2).exp()*epsilon
        return z

    def decode(self, x):
        return self.decoder(x)

    def forward(self, x):
        mean, logvar = self.encode(x)
        z = self.reparameterization(mean, logvar)
        x_hat = self.decode(z)
        return x_hat, mean, logvar
```

```
In [13]: model = VAE().to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```

```
In [14]: def loss_function(x, x_hat, mean, log_var):
    reproduction_loss = nn.functional.mse_loss(x_hat, x, reduction='sum')
    KLD = - 0.5 * torch.sum(1+ log_var - mean.pow(2) - log_var.exp())
```

```
return reproduction_loss + KLD
```

```
In [15]: def train(model, optimizer, epochs, device):
model.train()
for epoch in range(epochs):
    overall_loss = 0
    for batch_idx, (x, _) in enumerate(train_loader):
        x = x.view(batch_size, -1).to(device)

        optimizer.zero_grad()

        x_hat, mean, log_var = model(x)
        loss = loss_function(x, x_hat, mean, log_var)

        overall_loss += loss.item()

        loss.backward()
        optimizer.step()

    print("\tEpoch", epoch + 1, "\tAverage Loss: ", overall_loss/(batch_idx*batch_size))
return overall_loss

train(model, optimizer, epochs=50, device=device)
```



Epoch 1	Average Loss:	44.90734841679492
Epoch 2	Average Loss:	40.757268009345005
Epoch 3	Average Loss:	39.486435697679724
Epoch 4	Average Loss:	38.7257742233149
Epoch 5	Average Loss:	38.27964989663762
Epoch 6	Average Loss:	37.90731650325412
Epoch 7	Average Loss:	37.615150050702994
Epoch 8	Average Loss:	37.331490452022905
Epoch 9	Average Loss:	37.19050968329377
Epoch 10	Average Loss:	37.024223127412874
Epoch 11	Average Loss:	36.90037159914962
Epoch 12	Average Loss:	36.727448339191625
Epoch 13	Average Loss:	36.58070171069621
Epoch 14	Average Loss:	36.553968732066465
Epoch 15	Average Loss:	36.421423943062656
Epoch 16	Average Loss:	36.34540449903485
Epoch 17	Average Loss:	36.21433060227332
Epoch 18	Average Loss:	36.1550144079651
Epoch 19	Average Loss:	36.08352418011138
Epoch 20	Average Loss:	35.9948438722741
Epoch 21	Average Loss:	35.93922505119209
Epoch 22	Average Loss:	35.8742404210468
Epoch 23	Average Loss:	35.848164046196786
Epoch 24	Average Loss:	35.71230819269094
Epoch 25	Average Loss:	35.709727039369
Epoch 26	Average Loss:	35.68782412826717
Epoch 27	Average Loss:	35.61055909425866
Epoch 28	Average Loss:	35.57203330828073
Epoch 29	Average Loss:	35.51085223826821
Epoch 30	Average Loss:	35.470594671946735
Epoch 31	Average Loss:	35.42248072960142
Epoch 32	Average Loss:	35.40860261487244
Epoch 33	Average Loss:	35.31653475600611
Epoch 34	Average Loss:	35.29057205938935
Epoch 35	Average Loss:	35.28282936975037
Epoch 36	Average Loss:	35.24929938569491
Epoch 37	Average Loss:	35.216343539688545
Epoch 38	Average Loss:	35.16881060719689
Epoch 39	Average Loss:	35.121357393344375
Epoch 40	Average Loss:	35.10692839302483
Epoch 41	Average Loss:	35.096852032847714
Epoch 42	Average Loss:	35.08559472493218
Epoch 43	Average Loss:	34.998794373271856
Epoch 44	Average Loss:	34.98728417398137
Epoch 45	Average Loss:	35.02352690274807
Epoch 46	Average Loss:	34.88520460940761
Epoch 47	Average Loss:	34.962183530167465
Epoch 48	Average Loss:	34.86952581695404
Epoch 49	Average Loss:	34.86015463190604
Epoch 50	Average Loss:	34.8394119186274

```
Out[15]: 2086880.7739257812
```

```
In [16]: def generate_digit(z1, z2):
z_sample = torch.tensor([[z1, z2]], dtype=torch.float).to(device)
```

```
x_decoded = model.decode(z_sample)
digit = x_decoded.detach().cpu().reshape(28, 28) # reshape vector to 2d array
plt.figure(figsize = (2.5, 2.5))
plt.imshow(digit, cmap='gray')
plt.axis('off')
plt.show()

generate_digit(0.8, -0.9), generate_digit(-0.2, 1.0)
```



Out[16]: (None, None)

In []: